

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
# tiny synthetic DataFrame with Titanic-like columns
df = pd.DataFrame({
'survived': [0,1,1,0,1,0,1,0],
'pclass': [3,1,2,3,1,2,3,1],
'SEX': [' male',np.nan,'female','MaLe','female ','MALE','male','female'],
'age': [12,38,26,35,np.nan,54,2,27],
'sibsp': [1,3,0,1,0,0,3,0],
'parch': [0,1,0,0,0,0,1,2],
'fARE': [7.25,np.nan,7.925,8.05,53.1,51.8625,21.075,11.1333],
'embarcked': ['S','C','S','S','S','S',np.nan]
})
# Keep an original copy for comparisons later
df_orig = df.copy()
print(df.head())

```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarcked
0	0	3	male	12.0	1	0	7.250	S
1	1	1	NaN	38.0	3	1	NaN	C
2	1	2	female	26.0	0	0	7.925	S
3	0	3	MaLe	35.0	1	0	8.050	S
4	1	1	female	NaN	0	0	53.100	S

```

#1 Drop unnecessary columns from a DataFrame.
# Drop Unnecessary columns usinf drop()
df.drop(['embarked'], axis=1, inplace=True)
print(df.head()) # print

```

	survived	pclass	SEX	age	sibsp	parch	fARE
0	0	3	male	12.0	1	0	7.250
1	1	1	NaN	38.0	3	1	NaN
2	1	2	female	26.0	0	0	7.925
3	0	3	MaLe	35.0	1	0	8.050
4	1	1	female	NaN	0	0	53.100

```

#2 Rename columns for better readability.
# Rename pclass and sibsp using rename()
df.rename(columns={'pclass': 'Passenger_Class', 'sibsp':'Siblings_Spouses'}, inplace=True)
print(df.head()) # print

```

	survived	Passenger_Class	SEX	age	Siblings_Spouses	parch	fARE
0	0	3	male	12.0	1	0	7.250
1	1	1	NaN	38.0	3	1	NaN
2	1	2	female	26.0	0	0	7.925
3	0	3	MaLe	35.0	1	0	8.050
4	1	1	female	NaN	0	0	53.100

```

#3 Reorder columns in a DataFrame.
# Shuffle the column name based on order
df = df[['Passenger_Class', 'age', 'SEX', 'survived', 'Siblings_Spouses', 'parch', 'fARE']]
print(df.head()) # print

```

	Passenger_Class	age	SEX	survived	Siblings_Spouses	parch	fARE
0	3	12.0	male	0	1	0	7.250
1	1	38.0	NaN	1	3	1	NaN
2	2	26.0	female	1	0	0	7.925
3	3	35.0	MaLe	0	1	0	8.050
4	1	NaN	female	1	0	0	53.100

```

#4 Reset the index of a DataFrame
# Reset the index to original dataframe
df.reset_index(drop=True, inplace=True)
print(df.head()) # print

```

	Passenger_Class	age	SEX	survived	Siblings_Spouses	parch	fARE
0	3	12.0	male	0	1	0	7.250
1	1	38.0	NaN	1	3	1	NaN
2	2	26.0	female	1	0	0	7.925
3	3	35.0	MaLe	0	1	0	8.050
4	1	NaN	female	1	0	0	53.100

```
#5 Set a specific column as the new index
# Set index value as survived
df.set_index('survived', inplace=True)
print(df.head()) # print
```

	Passenger_Class	age	SEX	Siblings_Spouses	parch	fARE
survived						
0	3	12.0	male	1	0	7.250
1	1	38.0	NaN	3	1	NaN
1	2	26.0	female	0	0	7.925
0	3	35.0	MaLe	1	0	8.050
1	1	NaN	female	0	0	53.100

```
#6 Detect duplicate rows and remove them
# Remove duplicate rows using drop_duplicates()
df.drop_duplicates(inplace=True)
print(df.head()) # print
```

	Passenger_Class	age	SEX	Siblings_Spouses	parch	fARE
survived						
0	3	12.0	male	1	0	7.250
1	1	38.0	NaN	3	1	NaN
1	2	26.0	female	0	0	7.925
0	3	35.0	MaLe	1	0	8.050
1	1	NaN	female	0	0	53.100

```
#7 Detect inconsistent formatting in text columns and standardize them.
# Formate the inconsistent data
df['SEX'] = df['SEX'].str.strip()
print(df.head()) # print
```

	Passenger_Class	age	SEX	Siblings_Spouses	parch	fARE
survived						
0	3	12.0	male	1	0	7.250
1	1	38.0	NaN	3	1	NaN
1	2	26.0	female	0	0	7.925
0	3	35.0	male	1	0	8.050
1	1	NaN	female	0	0	53.100

```
#8 Convert all column names to lowercase
# convert column name to lowercase
df.columns = df.columns.str.lower()
print(df.head()) # print
```

	passenger_class	age	sex	siblings_spouses	parch	fare
survived						
0	3	12.0	male	1	0	7.250
1	1	38.0	NaN	3	1	NaN
1	2	26.0	female	0	0	7.925
0	3	35.0	male	1	0	8.050
1	1	NaN	female	0	0	53.100

```
#9 Identify missing values using .isnull().sum()
# Finding null value and compute sum
print(df.isnull().sum())
```

passenger_class	0
age	1
sex	1

```
---  
siblings_spouses    0  
parch                0  
fare                 1  
dtype: int64
```

```
#10 Drop rows with any missing values.  
# dropna() drops the missing value  
df.dropna(inplace=True)  
print(df.head()) # print
```

	survived	passenger_class	age	sex	siblings_spouses	parch	fare
0	0	3	12.0	male		1	7.2500
1	1	2	26.0	female		0	7.9250
0	0	3	35.0	male		1	8.0500
0	0	2	54.0	male		0	51.8625
1	1	3	2.0	male		3	21.0750

```
#11 Drop rows where a specific column has missing data  
#dropna() drops columun with missing data  
df.dropna(subset=['fARE'], inplace=True)  
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked
0	0	3	male	12.0	1	0	7.2500	S
2	1	2	female	26.0	0	0	7.9250	S
3	0	3	MaLe	35.0	1	0	8.0500	S
4	1	1	female	Nan	0	0	53.1000	S
5	0	2	MALE	54.0	0	0	51.8625	S

```
#12 Fill missing numeric values with column mean.  
# fill missing values with the mean value  
df['age'].fillna(df['age'].mean(), inplace=True)  
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked
0	0	3	male	12.0	1	0	7.2500	S
2	1	2	female	26.0	0	0	7.9250	S
3	0	3	MaLe	35.0	1	0	8.0500	S
4	1	1	female	26.0	0	0	53.1000	S
5	0	2	MALE	54.0	0	0	51.8625	S

/tmp/ipython-input-1128134902.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame. The behavior will change in pandas 3.0. This inplace method will never work because the int

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: v}

```
df['age'].fillna(df['age'].mean(), inplace=True)
```

```
#13 Fill missing categorical values with the mode.  
# fill missing value with mode value  
df['embarked'].fillna(df['embarked'].mode()[0], inplace=True)  
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked
0	0	3	male	12.0	1	0	7.250	S
1	1	1	NaN	38.0	3	1	Nan	C
2	1	2	female	26.0	0	0	7.925	S
3	0	3	MaLe	35.0	1	0	8.050	S
4	1	1	female	Nan	0	0	53.100	S

/tmp/ipython-input-2486643981.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame. The behavior will change in pandas 3.0. This inplace method will never work because the int

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: v}

```
df['embarked'].fillna(df['embarked'].mode()[0], inplace=True)
```

```
#14 Use forward fill and backward fill techniques.
```

```
# replace null with next value
df['age'].fillna(method='ffill')
# replace null with prev value
df['age'].fillna(method='bfill')
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked
0	0	3	male	12.0	1	0	7.250	S
1	1	1	Nan	38.0	3	1	Nan	C
2	1	2	female	26.0	0	0	7.925	S
3	0	3	MaLe	35.0	1	0	8.050	S
4	1	1	female	35.0	0	0	53.100	S

```
/tmp/ipython-input-4233193395.py:2: FutureWarning: Series.fillna with 'method' is deprecated
df['age'].fillna(method='ffill')
/tmp/ipython-input-4233193395.py:3: FutureWarning: Series.fillna with 'method' is deprecated
df['age'].fillna(method='bfill')
```

```
#15 Use interpolation to fill numeric gaps in time-series data
```

```
# Fill null values using interpolation
df['fARE'].interpolate(method='linear', direction='forward')
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked
0	0	3	male	12.0	1	0	7.2500	S
1	1	1	Nan	38.0	3	1	7.5875	C
2	1	2	female	26.0	0	0	7.9250	S
3	0	3	MaLe	35.0	1	0	8.0500	S
4	1	1	female	35.0	0	0	53.1000	S

```
#16 Create a new column as the sum of two existing columns.
```

```
# adds values of 2 columns
df['new_column'] = df['sibsp'] + df['parch']
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked	new_column
0	0	3	male	12.0	1	0	7.250	S	1
1	1	1	Nan	38.0	3	1	Nan	C	4
2	1	2	female	26.0	0	0	7.925	S	0
3	0	3	MaLe	35.0	1	0	8.050	S	1
4	1	1	female	Nan	0	0	53.100	S	0

```
#17 Create a binary column based on a condition.
```

```
# checks for age > 18
df['is_adult'] = df['age'] >= 18
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked	new_column \
0	0	3	male	12.0	1	0	7.250	S	1
1	1	1	Nan	38.0	3	1	Nan	C	4
2	1	2	female	26.0	0	0	7.925	S	0
3	0	3	MaLe	35.0	1	0	8.050	S	1
4	1	1	female	Nan	0	0	53.100	S	0

```
is_adult
0    False
1     True
2     True
3     True
4    False
```

```
#18 Extract year/month/day from a datetime column.
```

```
# range() is used to generate date
df['date'] = pd.date_range('2022-01-01', periods=len(df), freq='D')
# prints year
df['year'] = df['date'].dt.year
..
```

```
# prints month
df['month'] = df['date'].dt.month
# prints date
df['day'] = df['date'].dt.day
# print dataset
print(df.head())
```

	age	fare	embarked	survived	date	year	month	day
0	22	7.25	S	0	2022-01-01	2022	1	1
1	38	71.83	C	1	2022-01-02	2022	1	2
2	26	8.05	S	1	2022-01-03	2022	1	3
3	35	53.10	S	1	2022-01-04	2022	1	4
4	28	8.46	S	0	2022-01-05	2022	1	5

```
#19 Create a column showing length of string in a text field
# prints length of the data present in sex
df['sex_length'] = df['SEX'].str.len()
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked	sex_length
0	0	3	male	12.0	1	0	7.250	S	5.0
1	1	1	NaN	38.0	3	1	NaN	C	NaN
2	1	2	female	26.0	0	0	7.925	S	6.0
3	0	3	MaLe	35.0	1	0	8.050	S	4.0
4	1	1	female	NaN	0	0	53.100	S	7.0

```
#20 Bin a continuous variable into fixed intervals
# bin age value into discrete intervals
bins = [0, 18, 65, float('inf')]
# add labels based on values
labels = ['minor', 'adult', 'senior']
df['age_group'] = pd.cut(df['age'], bins=bins, labels=labels)
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked	sex_length
0	0	3	male	12.0	1	0	7.250	S	5.0
1	1	1	NaN	38.0	3	1	NaN	C	NaN
2	1	2	female	26.0	0	0	7.925	S	6.0
3	0	3	MaLe	35.0	1	0	8.050	S	4.0
4	1	1	female	NaN	0	0	53.100	S	7.0

  

	age_group
0	minor
1	adult
2	adult
3	adult
4	NaN

```
#21 Create a new column showing whether a value is above or below median
# checks for fare whether above or below median
df['fare_above_median'] = df['fARE'] > df['fARE'].median()
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	fARE	embarked	sex_length
0	0	3	male	12.0	1	0	7.250	S	5.0
1	1	1	NaN	38.0	3	1	NaN	C	NaN
2	1	2	female	26.0	0	0	7.925	S	6.0
3	0	3	MaLe	35.0	1	0	8.050	S	4.0
4	1	1	female	NaN	0	0	53.100	S	7.0

  

	age_group	fare_above_median
0	minor	False
1	adult	False
2	adult	False
3	adult	False
4	NaN	True

```
#22 Encode a categorical variable using one-hot encoding
```

```
# converts categorial data into a numerical format
df = pd.get_dummies(df, columns=['embarked'])
print(df.head()) # print
```

	survived	pclass	SEX	age	sibsp	parch	FARE	sex_length
0	0	3	male	12.0	1	0	7.250	5.0
1	1	1	NaN	38.0	3	1	NaN	NaN
2	1	2	female	26.0	0	0	7.925	6.0
3	0	3	MaLe	35.0	1	0	8.050	4.0
4	1	1	female	NaN	0	0	53.100	7.0

  

	age_group	fare_above_median	embarked_C	embarked_S
0	minor	False	False	True
1	adult	False	True	False
2	adult	False	False	True
3	adult	False	False	True
4	NaN	True	False	True

```
# Use the following for problem # 23
# Sample dataset
data = {
    'age': [22, 38, 26, 35, 28, 54, 2, 19, 40, 30],
    'fare': [7.25, 71.83, 8.05, 53.10, 8.46, 51.86, 21.07, 30.0, 27.75, 13.0],
    'embarked': ['S', 'C', 'S', 'S', 'S', 'Q', 'S', 'C', 'Q', 'S'],
    'survived': [0, 1, 1, 1, 0, 0, 1, 0, 1, 0] # target variable
}
df = pd.DataFrame(data)
# Features and target
x = df[['age', 'fare', 'embarked']]
y = df['survived']
print(df.head())
```

	age	fare	embarked	survived
0	22	7.25	S	0
1	38	71.83	C	1
2	26	8.05	S	1
3	35	53.10	S	1
4	28	8.46	S	0

```
#23 Split data into train and test data
# Importing the train_test_split function from sklearn for splitting the dataset
from sklearn.model_selection import train_test_split
# Splitting the dataset into training (80%) and testing (20%) sets
# 'x' is the feature data, and 'y' is the target (labels) data
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
# Display the first few rows of the training feature data
print(x_train.head())
# Display the first few rows of the testing feature data
print(x_test.head())
# Display the first few rows of the training target labels
print(y_train.head())
# Display the first few rows of the testing target labels
print(y_test.head())
```

	age	fare	embarked
5	54	51.86	Q
0	22	7.25	S
7	19	30.00	C
2	26	8.05	S
9	30	13.00	S

  

	age	fare	embarked
8	40	27.75	Q
1	38	71.83	C
5	0		
0	0		
7	0		
2	1		
9	0		

```
Name: survived, dtype: int64  
8      1  
1      1  
Name: survived, dtype: int64
```